


# An infinitely customizable training loop

Sylvain Gugger



The basic training loop iterates repeatedly over 4 simple steps.



And in PyTorch,  
this fits in just  
five lines of  
codes.

```
def train(train_dl, model, epoch, opt, loss_func):  
    for _ in range(epoch):  
        model.train()  
        for xb, yb in train_dl:  
            out = model(xb)  
            loss = loss_func(out, yb)  
            loss.backward()  
            opt.step()  
            opt.zero_grad()
```

But there are multiple  
tweaks you can add to it.

# Keeping track of losses and metrics




```
to_fp16(learn)

%time fit_one_cycle(learn, 3e-3, 30, wd=0.2)
```

0.0% [0/30 00:00<00:00]



# Hyper- parameters schedules



# Every regularization technique

## Averaging Weights Leads to Wider Optima and Better Generalization

Pavel Izmailov<sup>\*1</sup> Dmitrii Podoprikin<sup>\*2,3</sup> Timur Garipov<sup>\*4,5</sup> Dmitry Vetrov<sup>2,3</sup> Andrew Gordon Wilson<sup>1</sup>  
<sup>1</sup>Cornell University, <sup>2</sup>Higher School of Economics, <sup>3</sup>Samsung-HSE Laboratory,  
<sup>4</sup>Samsung AI Center in Moscow, <sup>5</sup>Lomonosov Moscow State University




Figure 3: The  $L_2$ -regularized cross-entropy train loss and test error surfaces of a Preactivation ResNet-164 on CIFAR-100 in the plane containing the first, middle and last points (indicated by black crosses) in the trajectories with (left two) cyclical and (right two) constant learning rate schedules.

### 2.1. $L_2$ activation regularization (AR)

While  $L_2$  regularization is traditionally used on the weights of machine learning models ( $L_2$  weight decay), it could also be used on the activations. We define AR as

$$\alpha L_2(m \odot h_t)$$

where  $m$  is the dropout mask used by later parts of the model,  $L_2(\cdot) = \|\cdot\|_2$  ( $L_2$  norm),  $h_t$  is the output of the RNN at timestep  $t$ , and  $\alpha$  is a scaling coefficient.

Temporal activation regularization (TAR) is a direct descendant of this slowness regularization, minimizing

$$\beta L_2(h_t - h_{t+1})$$

where  $L_2(\cdot) = \|\cdot\|_2$  ( $L_2$  norm),  $h_t$  is the output of the RNN at timestep  $t$ , and  $\beta$  is a scaling coefficient.

## DECOUPLED WEIGHT DECAY REGULARIZATION

Ilya Loshchilov & Frank Hutter

University of Freiburg  
 Freiburg, Germany,  
 {ilya, fh}@cs.uni-freiburg.de

### Algorithm 2 Adam with $L_2$ regularization and Adam with decoupled weight decay (AdamW)

```

1: given  $\alpha = 0.001, \beta_1 = 0.9, \beta_2 = 0.999, \epsilon = 10^{-8}, \lambda \in \mathbb{R}$ 
2: initialize time step  $t \leftarrow 0$ , parameter vector  $\theta_{t=0} \in \mathbb{R}^n$ , first moment vector  $m_{t=0} \leftarrow \mathbf{0}$ , second moment vector  $v_{t=0} \leftarrow \mathbf{0}$ , schedule multiplier  $\eta_{t=0} \in \mathbb{R}$ 
3: repeat
4:    $t \leftarrow t + 1$ 
5:    $\nabla f_t(\theta_{t-1}) \leftarrow \text{SelectBatch}(\theta_{t-1})$  ▷ select batch and return the corresponding gradient
6:    $g_t \leftarrow \nabla f_t(\theta_{t-1}) + \lambda \theta_{t-1}$ 
7:    $m_t \leftarrow \beta_1 m_{t-1} + (1 - \beta_1) g_t$  ▷ here and below all operations are element-wise
8:    $v_t \leftarrow \beta_2 v_{t-1} + (1 - \beta_2) g_t^2$ 
9:    $\hat{m}_t \leftarrow m_t / (1 - \beta_1^t)$  ▷  $\beta_1$  is taken to the power of  $t$ 
10:   $\hat{v}_t \leftarrow v_t / (1 - \beta_2^t)$  ▷  $\beta_2$  is taken to the power of  $t$ 
11:   $\eta_t \leftarrow \text{SetScheduleMultiplier}(t)$  ▷ can be fixed, decay, or also be used for warm restarts
12:   $\theta_t \leftarrow \theta_{t-1} - \eta_t \left( \alpha \hat{m}_t / (\sqrt{\hat{v}_t} + \epsilon) + \lambda \theta_{t-1} \right)$ 
13: until stopping criterion is met
14: return optimized parameters  $\theta_t$ 
    
```

## *mixup*: BEYOND EMPIRICAL RISK MINIMIZATION

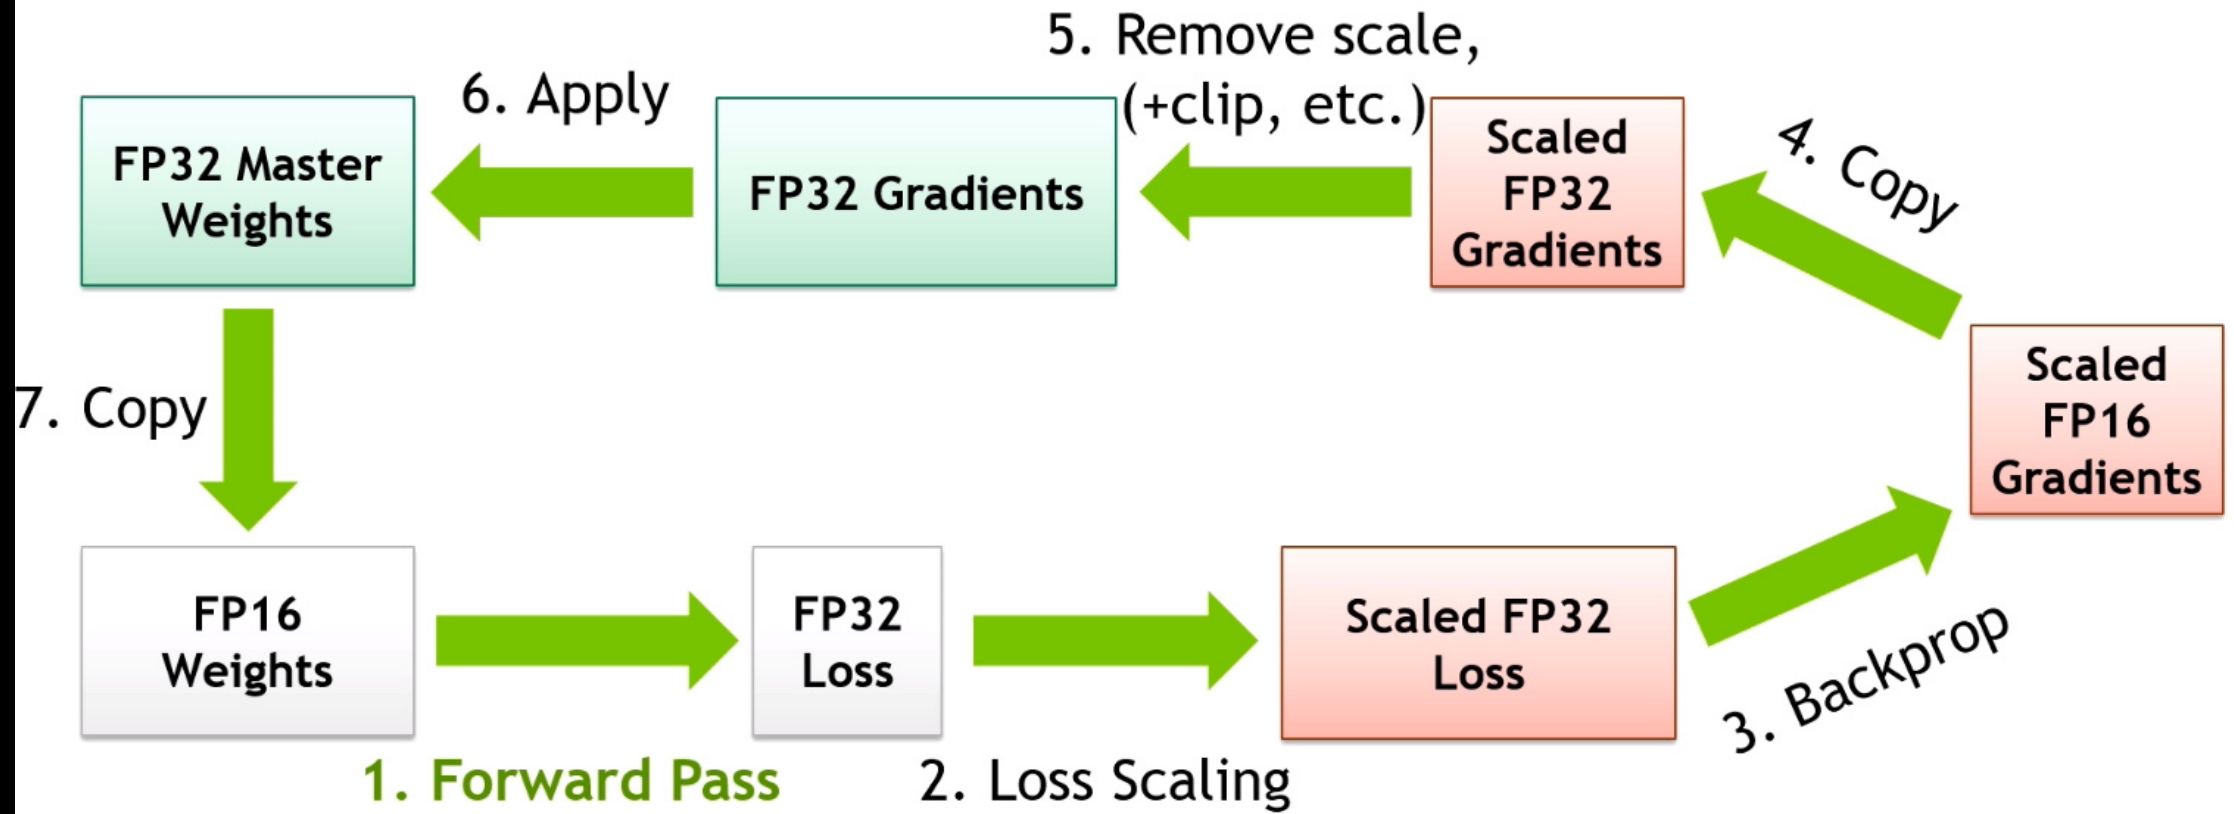
Hongyi Zhang  
 MIT

Moustapha Cisse, Yann N. Dauphin, David Lopez-Paz<sup>\*</sup>  
 FAIR

**Contribution** Motivated by these issues, we introduce a simple and data-agnostic data augmentation routine, termed *mixup* (Section 2). In a nutshell, *mixup* constructs virtual training examples


$$\begin{aligned} \tilde{x} &= \lambda x_i + (1 - \lambda) x_j, & \text{where } x_i, x_j \text{ are raw input vectors} \\ \tilde{y} &= \lambda y_i + (1 - \lambda) y_j, & \text{where } y_i, y_j \text{ are one-hot label encodings} \end{aligned}$$

# MIXED PRECISION TRAINING





Or more  
complex  
trainings like  
GANs



So you either  
have to rewrite a  
new loop for  
each new kind of  
training...



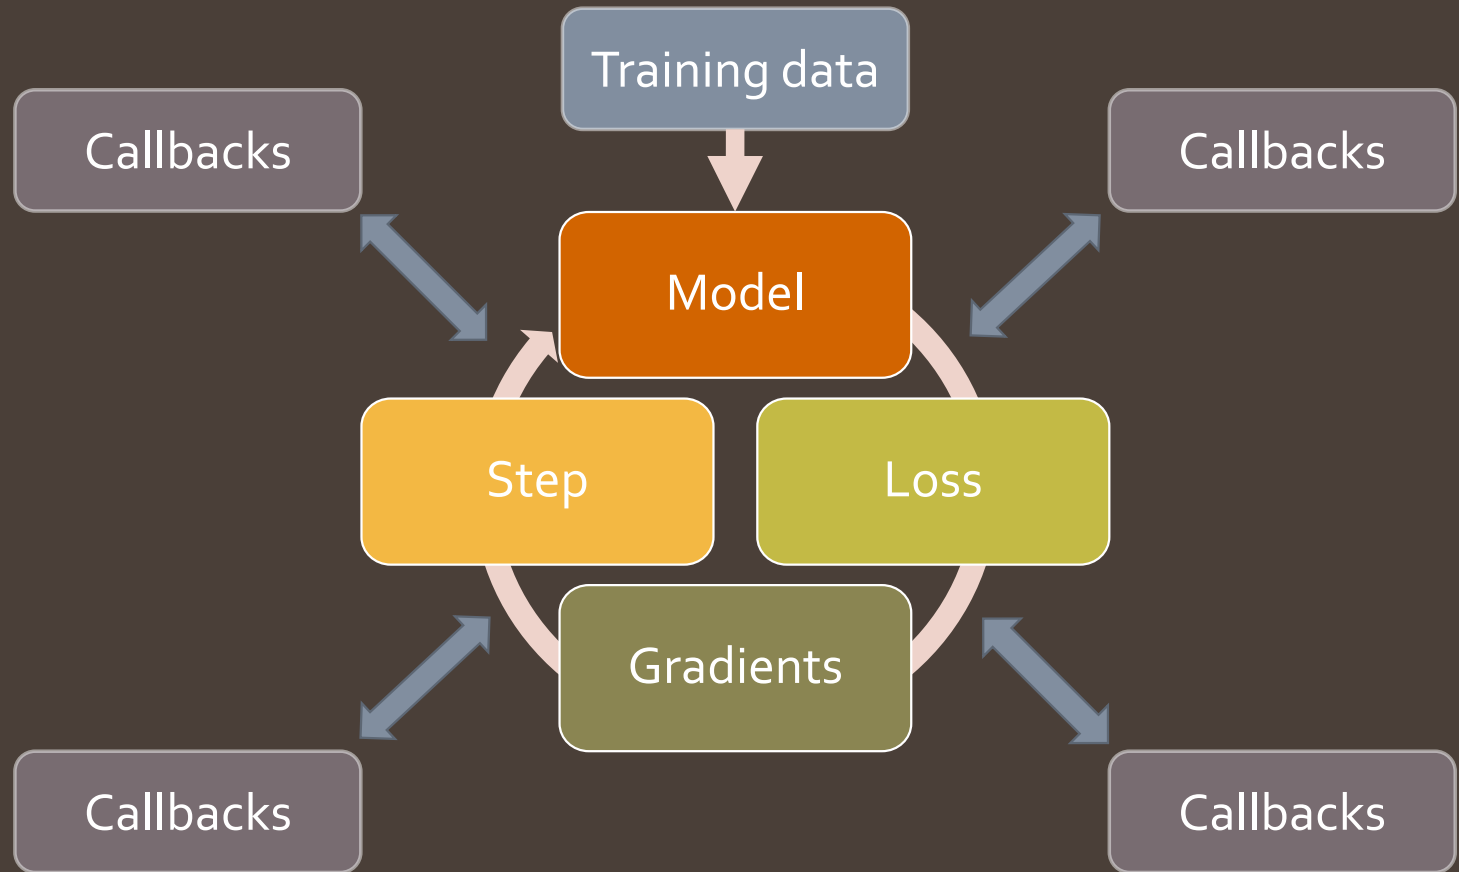
Or try to write something that incorporates everything you can think of.

```
class Stepper():
    def __init__(self, m, opt, crit, clip=0, reg_fn=None, fp16=False, loss_scale=1):
        self.m, self.opt, self.crit, self.clip, self.reg_fn = m, opt, crit, clip, reg_fn
        self.fp16 = fp16
        self.reset(True)
        if self.fp16: self.fp32_params = copy_model_to_fp32(m, opt)
        self.loss_scale = loss_scale

    def reset(self, train=True):
        if train: apply_leaf(self.m, set_train_mode)
        else: self.m.eval()
        if hasattr(self.m, 'reset'):
            self.m.reset()
            if self.fp16: self.fp32_params = copy_model_to_fp32(self.m, self.opt)

    def step(self, xs, y, epoch):
        xtra = []
        output = self.m(*xs)
        if isinstance(output, tuple): output, *xtra = output
        if self.fp16: self.m.zero_grad()
```

Fortunately there is  
a way around this:  
Callbacks




Going back  
from our basic  
training loop...

```
def train(train_dl, model, epoch, opt, loss_func):  
    for _ in range(epoch):  
        model.train()  
        for xb, yb in train_dl:  
            out = model(xb)  
            loss = loss_func(out, yb)  
            loss.backward()  
            opt.step()  
            opt.zero_grad()
```

...we can just  
add callbacks  
everywhere.

```
def train(train_dl, model, epoch, opt, loss_func, callbacks):
    callbacks.on_train_begin()
    for _ in range(epoch):
        callbacks.on_epoch_begin()
        model.train()
        for xb, yb in train_dl:
            callbacks.on_batch_begin()
            out = model(xb)
            callbacks.on_loss_begin()
            loss = loss_func(out, yb)
            callbacks.on_loss_end()
            loss.backward()
            callbacks.on_step_begin()
            opt.step()
            callbacks.on_step_end()
            opt.zero_grad()
            callbacks.on_batch_end()
        callbacks.on_epoch_end()
    callbacks.on_train_end()
```






But for this to  
be infinitely  
flexible, each  
callback need to  
see everything  
that's going  
on...


... and be able to change  
anything in the training loop.

So we add a handler that stores all the information in the training loop to feed to each callback.





Which in turn can all update the state of the handler, which it gives back to the training loop.



The general context is provided by grouping the model, data, loss function and optimizer in a Learner class.

```
@dataclass
class DataBunch():
    train_dl:DataLoader
    valid_dl:DataLoader
```

```
@dataclass
class BasicLearner():
    model:nn.Module
    loss_func:LossFunction
    opt:optim.Optimizer
    data:DataBunch
```



Training loop

Those updates  
can be new  
values, or flags  
that skip steps  
or stop the  
training.

```
def train(learn, epochs, callbacks, metrics):
    cb_handler = CallbackHandler(callbacks)
    cb_handler.on_train_begin(epochs, learn, metrics)
    for epoch in range(epochs):
        learn.model.train()
        cb_handler.on_epoch_begin(epoch)
        for xb,yb in learn.data.train_dl:
            xb,yb = cb_handler.on_batch_begin(xb,yb)
            out = learn.model(xb)
            out = cb_handler.on_loss_begin(out)
            loss = learn.loss_func(out, yb)
            loss, skip_backward = callbacks.on_loss_begin(loss)
            if not skip_backward: loss.backward()
            if not cb_handler.on_step_begin(): learn.opt.step()
            if not cb_handler.on_step_end(): learn.opt.zero_grad()
            if not callbacks.on_batch_end(): break
        val_loss, mets = validate(learn.data.valid_dl, model, metrics)
        if not callbacks.on_epoch_end(val_loss, mets): break
    callbacks.on_train_end()
```

Then each tweak of the training loop can be entirely written in its own callback.

```
class LRScheduler(LearnerCallback):
```

```
    def on_batch_begin(self, iteration, **kwargs):  
        self.learn.opt.lr = my_function(iteration)
```

```
class EarlyStopping(Callback):
```

```
    def on_epoch_end(self, last_metrics, **kwargs):  
        if not_happy(last_metrics): return {'stop_training': True}
```

```
class ParallelTrainer(LearnerCallback):
```

```
    _order = -20  
    def on_train_begin(self, **kwargs):  
        self.learn.model = DataParallel(self.learn.model)  
    def on_train_end(self, **kwargs):  
        self.learn.model = self.learn.model.module
```

```
class GradientClipping(LearnerCallback):
```

```
    "Gradient clipping during training."
```

```
    def __init__(self, learn:Learner, clip:float = 0.):  
        super().__init__(learn)  
        self.clip = clip
```

```
    def on_backward_end(self, **kwargs):
```

```
        "Clip the gradient before the optimizer step."
```

```
        if self.clip: nn.utils.clip_grad_norm_(self.learn.model.parameters(), self.clip)
```

Even if they  
are a bit more  
complex.

```
class AccumulateScheduler(LearnerCallback):
    "Does accumulated step every nth step by accumulating gradients"
    def __init__(self, learn:Learner, n_step:int = 1, drop_last:bool = False):
        super().__init__(learn)
        self.n_step, self.drop_last = n_step, drop_last
        if hasattr(self.loss_func, "reduction") and (self.loss_func.reduction != "sum"):
            warn("For better gradients consider 'reduction=sum'")

    def on_epoch_begin(self, **kwargs):
        "Init samples and batches."
        self.acc_samples, self.acc_batches = 0., 0.

    def on_batch_begin(self, last_input, last_target, **kwargs):
        "Accumulate samples and batches"
        self.acc_samples += last_input.shape[0]
        self.acc_batches += 1

    def on_backward_end(self, **kwargs):
        "Accumulated step and reset samples."
        if (self.acc_batches % self.n_step) == 0:
            for p in (self.learn.model.parameters()):
                if p.requires_grad: p.grad.div_(self.acc_samples)
            self.acc_samples = 0
        else: return {'skip_step':True, 'skip_zero':True}

    def on_epoch_end(self, **kwargs):
        "Step the rest of the accumulated grads if not perfectly divisible"
        for p in (self.learn.model.parameters()):
            if p.requires_grad: p.grad.div_(self.acc_samples)
        if not self.drop_last: self.learn.opt.step()
        self.learn.opt.zero_grad()
```

It's then easy to  
mix and match  
those blocks  
together...

...and to add a  
new one!

SGDR

Record  
metrics

Multi-GPU

Mixed  
Precision

GAN

Gradient  
clipping

1cycle

SWA

Mixup

Save best  
model

Tensoboard

Early  
stopping

Gradient  
accumulation

Loss  
penalties

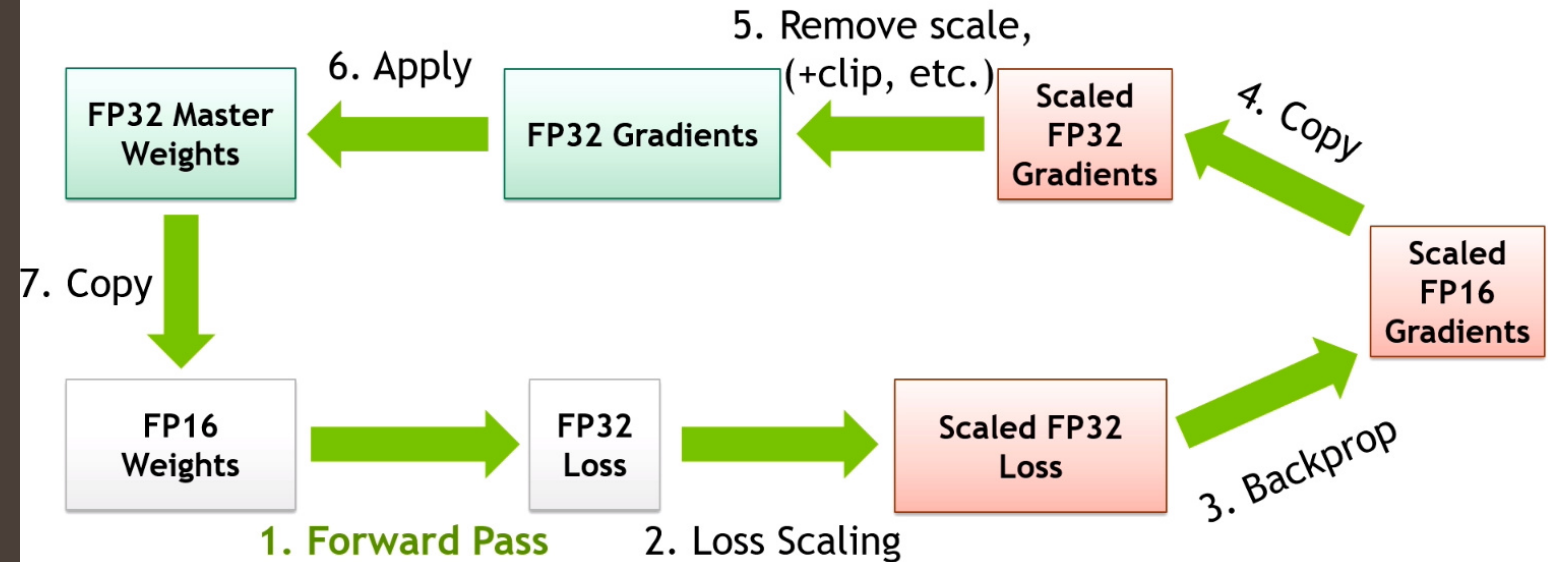
AdamW

Yours!


```
learn.fit(epochs, lr, wd, callbacks)
```

# Case study: Mixed precision training

## MIXED PRECISION TRAINING



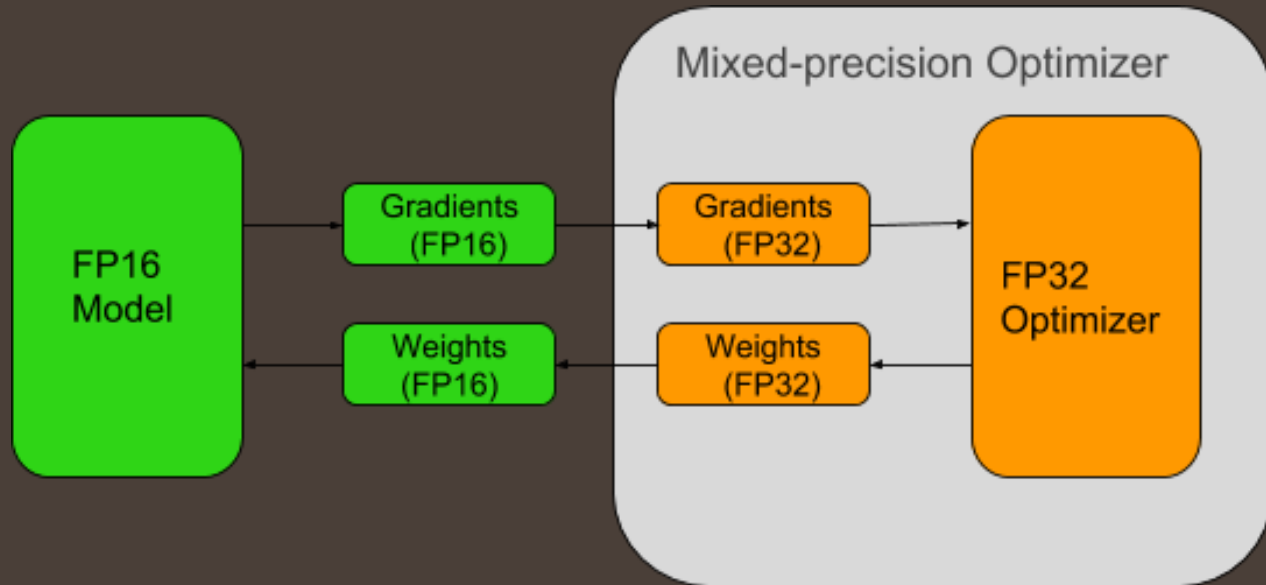
## HALF-PRECISION FLOAT (FLOAT16)



FLOAT16 has wide range ( $2^{40}$ ) ... but not as wide as FP32!

Normal range:  $[6 \times 10^{-5}, 65504]$   
Sub-normal range:  $[6 \times 10^{-8}, 6 \times 10^{-5}]$

We define master parameters (in FP32) and model parameters (in FP16)

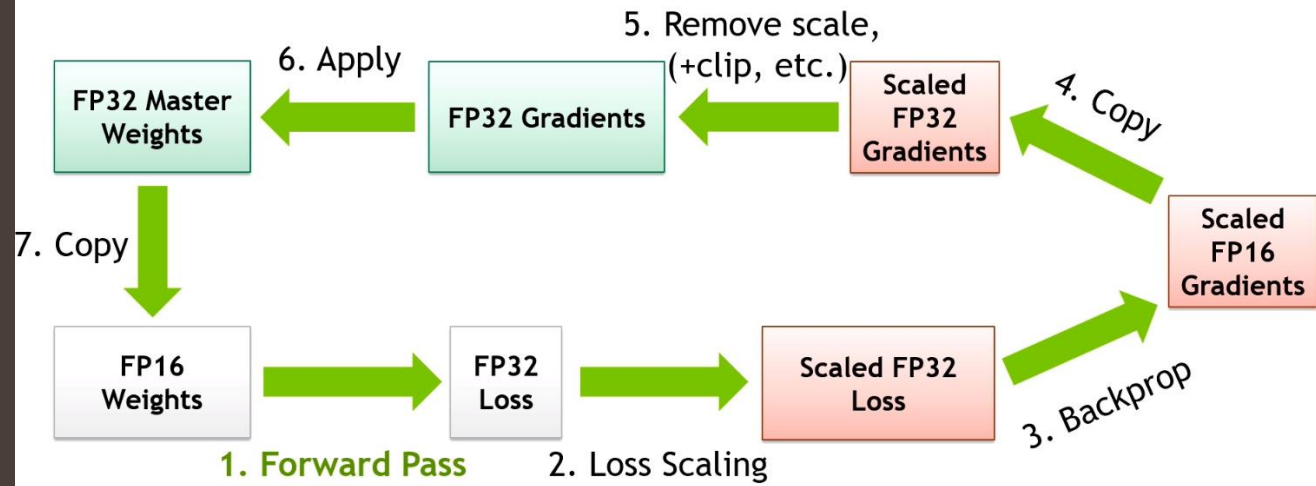


```
class MixedPrecision(LearnerCallback):
    _order = 999 #Need to run after things that could call on_backward_begin and change the loss
    "Callback that handles mixed-precision training."
    def __init__(self, learn:Learner, loss_scale:float=None, clip:float=None,
                 flat_master:bool=False):
        super().__init__(learn)
        self.flat_master, self.loss_scale, self.clip = flat_master, loss_scale, clip
        assert torch.backends.cudnn.enabled, "Mixed precision training requires cudnn."

    def on_train_begin(self, **kwargs:Any)->None:
        "Prepare the master model."
        #Get a copy of the model params in FP32
        self.model_params, self.master_params = get_master(self.learn.layer_groups, self.flat_master)
        #Changes the optimizer so that the optimization step is done in FP32.
        self.learn.opt = self.learn.opt.new_with_params(self.master_params)
```



And we can easily add the relevant parts in the appropriate event function.



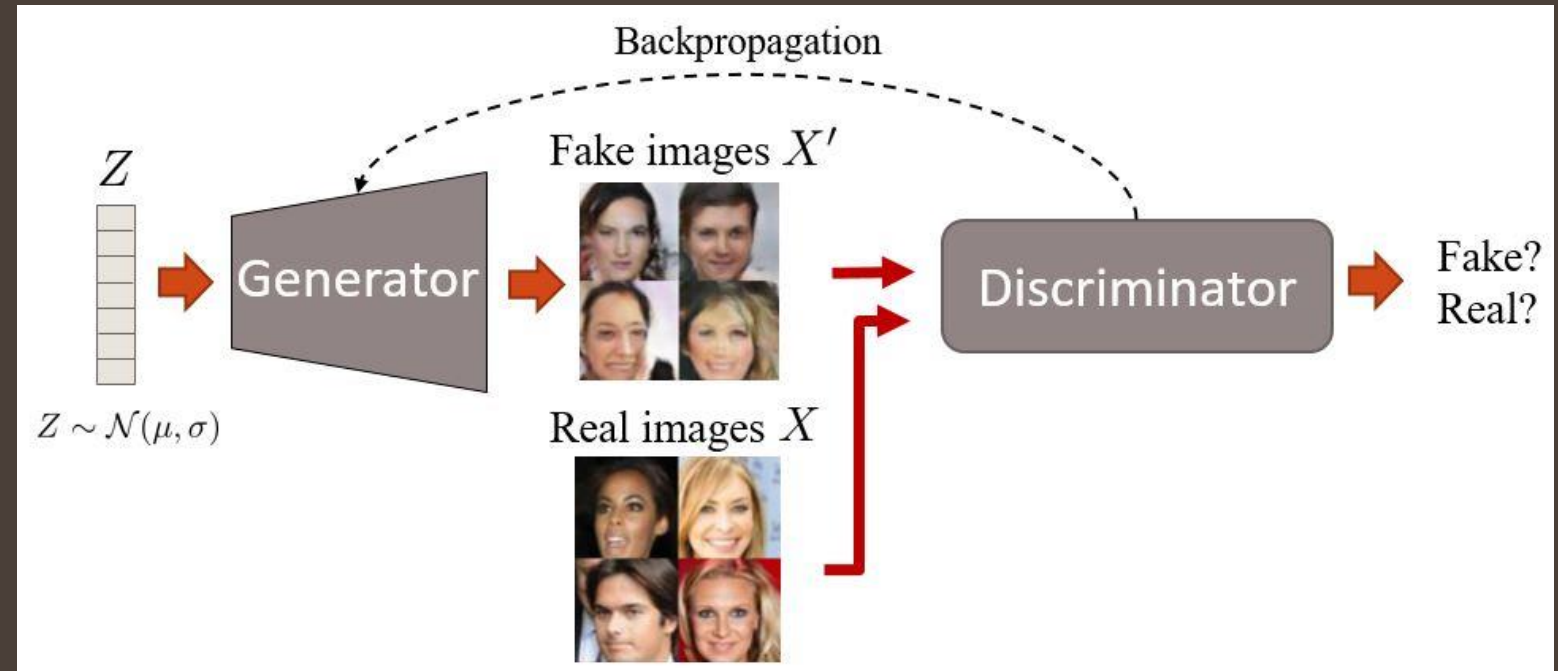
```
def on_loss_begin(self, last_output:Tensor, **kwargs:Any) -> Tensor:
    "Convert half precision output to FP32 to avoid reduction overflow."
    return {'last_output': to_float(last_output)}
```

```
def on_backward_begin(self, last_loss:Rank0Tensor, **kwargs:Any) -> Rank0Tensor:
    "Scale gradients up by `self.loss_scale` to prevent underflow."
    #To avoid gradient underflow, we scale the gradients
    ret_loss = last_loss * self.loss_scale
    return {'last_loss': ret_loss}
```

```
def on_backward_end(self, **kwargs:Any)->None:
    "Convert the gradients back to FP32 and divide them by the scale."
    model_g2master_g(self.model_params, self.master_params, self.flat_master)
    for group in self.master_params:
        for param in group:
            if param.grad is not None: param.grad.div_(self.loss_scale)
    if self.clip is not None:
        for group in self.master_params: nn.utils.clip_grad_norm_(group, self.clip)
```

```
def on_step_end(self, **kwargs:Any)->None:
    "Update the params from master to model and zero grad."
    #Zeros the gradients of the model since the optimizer is disconnected.
    self.learn.model.zero_grad()
    master2model(self.model_params, self.master_params, self.flat_master)
```

# Case study: GAN Training



We define a model and a loss function that wraps critic and generator with an inner switch.

```
class GANModule(nn.Module):
    "Wrapper around a `generator` and a `critic` to create a GAN."
    def __init__(self, generator:nn.Module=None, critic:nn.Module=None, gen_mode:bool=False):
        super().__init__()
        self.gen_mode = gen_mode
        if generator: self.generator,self.critic = generator,critic

    def forward(self, *args):
        return self.generator(*args) if self.gen_mode else self.critic(*args)

    def switch(self, gen_mode:bool=None):
        "Put the model in generator mode if `gen_mode`, in critic mode otherwise."
        self.gen_mode = (not self.gen_mode) if gen_mode is None else gen_mode
```

```
class GANLoss(GANModule):
    "Wrapper around `loss_funcC` (for the critic) and `loss_funcG` (for the generator)."
```

```
    def __init__(self, loss_funcG:Callable, loss_funcC:Callable, gan_model:GANModule):
        super().__init__()
        self.loss_funcG,self.loss_funcC,self.gan_model = loss_funcG,loss_funcC,gan_model

    def generator(self, output, target):
        """Evaluate the `output` with the critic then uses `self.loss_funcG` to combine
        it with `target`."""
        fake_pred = self.gan_model.critic(output)
        return self.loss_funcG(fake_pred, target, output)

    def critic(self, real_pred, input):
        """Create some `fake_pred` with the generator from `input` and compare them to `real_pred`
        in `self.loss_funcD`."""
        fake = self.gan_model.generator(input.requires_grad_(False)).requires_grad_(True)
        fake_pred = self.gan_model.critic(fake)
        return self.loss_funcC(real_pred, fake_pred)
```

Then the base of the GANTrainer Callback is to switch from the generator to the critic.

```
class GANTrainer(LearnerCallback):
    "Handles GAN Training."
    _order=-20
    def __init__(self, learn:Learner, clip:float=None, gen_first:bool=False):
        super().__init__(learn)
        self.clip,self.gen_first = clip,gen_first
        self.generator,self.critic = self.model.generator,self.model.critic

    def _set_trainable(self):
        train_model = self.generator if self.gen_mode else self.critic
        loss_model = self.generator if not self.gen_mode else self.critic
        requires_grad(train_model, True)
        requires_grad(loss_model, False)

    def switch(self, gen_mode:bool=None):
        "Switch the model, if `gen_mode` is provided, in the desired mode."
        self.gen_mode = (not self.gen_mode) if gen_mode is None else gen_mode
        self.opt = self.opt_gen if self.gen_mode else self.opt_critic
        self._set_trainable()
        self.model.switch(gen_mode)
        self.loss_func.switch(gen_mode)
```

Then it's just a matter of making sure we have things fed to the model according to the right mode.

```
def on_train_begin(self, **kwargs):
    "Create the optimizers for the generator and critic if necessary, initialize smootheners."
    if not getattr(self, 'opt_gen', None):
        self.opt_gen = self.opt.new([nn.Sequential(*flatten_model(self.generator))])
    else: self.opt_gen.lr, self.opt_gen.wd = self.opt.lr, self.opt.wd
    if not getattr(self, 'opt_critic', None):
        self.opt_critic = self.opt.new([nn.Sequential(*flatten_model(self.critic))])
    else: self.opt_critic.lr, self.opt_critic.wd = self.opt.lr, self.opt.wd
    self.gen_mode = self.gen_first
    self.switch(self.gen_mode)
    self.closses, self.glosses = [], []
    self.recorder.add_metric_names(['gen_loss', 'disc_loss'])
```

```
def on_train_end(self, **kwargs):
    "Switch in generator mode for showing results."
    self.switch(gen_mode=True)
```

```
def on_batch_begin(self, last_input, last_target, **kwargs):
    "Clamp the weights with `self.clip` if it's not None, return the correct input."
    if self.clip is not None:
        for p in self.critic.parameters(): p.data.clamp_(-self.clip, self.clip)
    if self.gen_mode: return {'last_input': last_input, 'last_target': last_target}
    else: return {'last_input': last_target, 'last_target': last_input}
```

```
def on_backward_begin(self, last_loss, last_output, **kwargs):
    "Record `last_loss` in the proper list."
    if self.gen_mode: self.glosses.append(last_loss.detach().cpu())
    else: self.closses.append(last_loss.detach().cpu())
```

```
def on_epoch_begin(self, epoch, **kwargs):
    "Put the critic or the generator back to eval if necessary."
    self.switch(self.gen_mode)
```

```
def on_epoch_end(self, pbar, epoch, last_metrics, **kwargs):
    "Put the various losses in the recorder and show a sample image."
    return add_metrics(last_metrics, self.glosses[-1], self.closses[-1])
```

Then another  
Callback is  
responsible for  
switching back  
and forth

```
class FixedGANSwitcher(LearnerCallback):
    "Switcher to do `n_crit` iterations of the critic then `n_gen` iterations of the generator."
    def __init__(self, learn:Learner, n_crit:Union[int,Callable]=1, n_gen:Union[int,Callable]=1):
        super().__init__(learn)
        self.n_crit,self.n_gen = n_crit,n_gen

    def on_train_begin(self, **kwargs):
        "Initiate the iteration counts."
        self.n_c,self.n_g = 0,0

    def on_batch_end(self, iteration, **kwargs):
        "Switch the model if necessary."
        if self.learn.gan_trainer.gen_mode:
            self.n_g += 1
            n_iter,n_in,n_out = self.n_gen,self.n_c,self.n_g
        else:
            self.n_c += 1
            n_iter,n_in,n_out = self.n_crit,self.n_g,self.n_c
        target = n_iter if isinstance(n_iter, int) else n_iter(n_in)
        if target == n_out:
            self.learn.gan_trainer.switch()
            self.n_c,self.n_g = 0,0
```

# In conclusion

A well-designed Callback system allows you to:

- keep the training loop as simple as possible
- have each tweak independently written in its own Callback
- easily mix and match, or perform ablation studies
- easily add a new experiment idea and debug it
- make it simple for contributors to add their own

You can use this system and all our Callbacks by installing fastai:

```
conda install -c pytorch -c fastai fastai
```